Coursera Deep Learning Course 2-Week 2: Optimization Methods

Kang Zhao

2018.09.06

Course Contents

- Mini-batch gradient descent
- Gradient Descent with momentum
- RMSprop
- Adam
- Learning Rate Decay

Optimization Overview

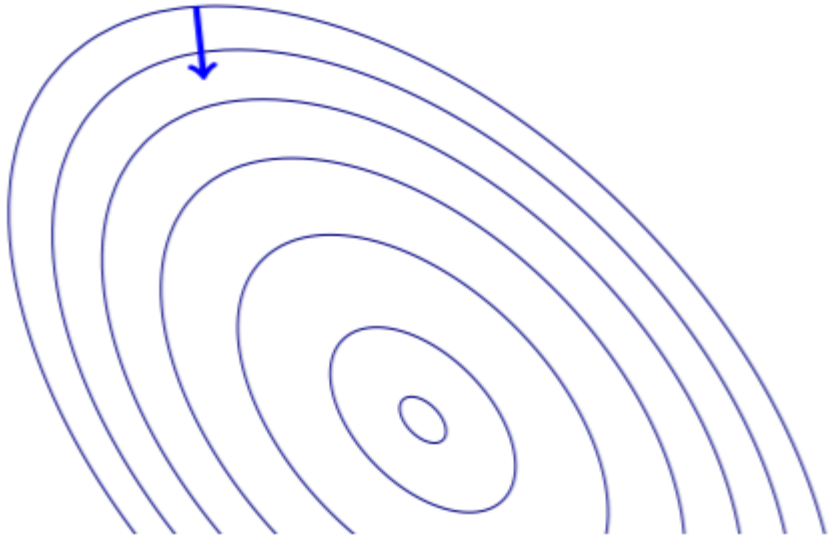

Gradient Descent

Mini-Batch & SGD

- Stochastic Gradient Descent
  - Train only one sample at one epoch
- Batch Gradient Descent
  - Train all the samples at one epoch
- Mini Batch Gradient Descent
  - Train k samples together at one epoch

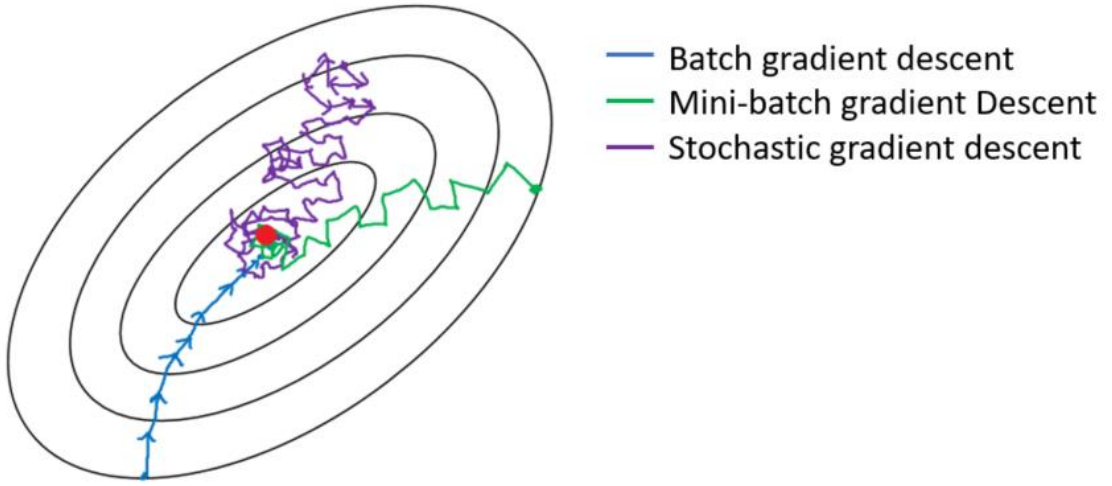

Mini-batch size: K

Exponentially Weighted (Moving) Average

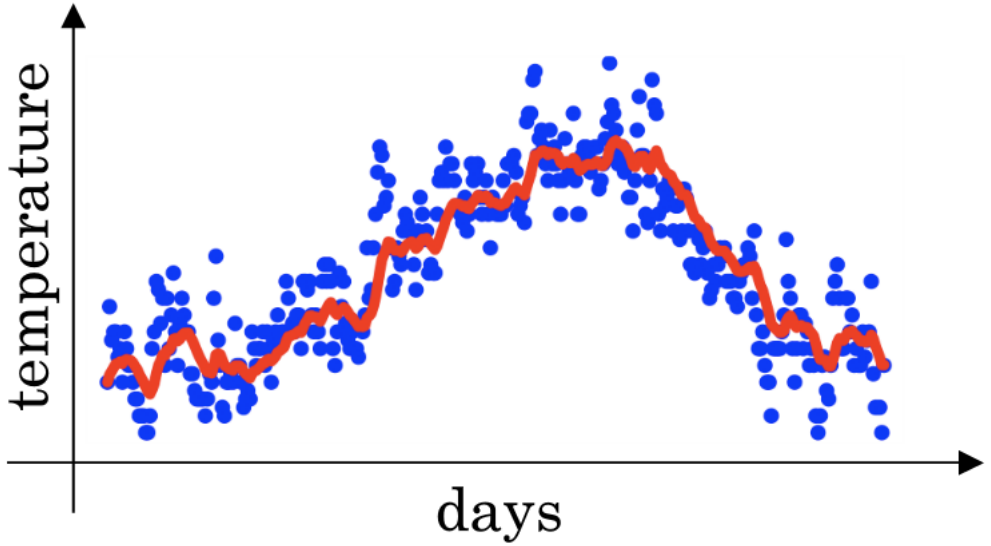

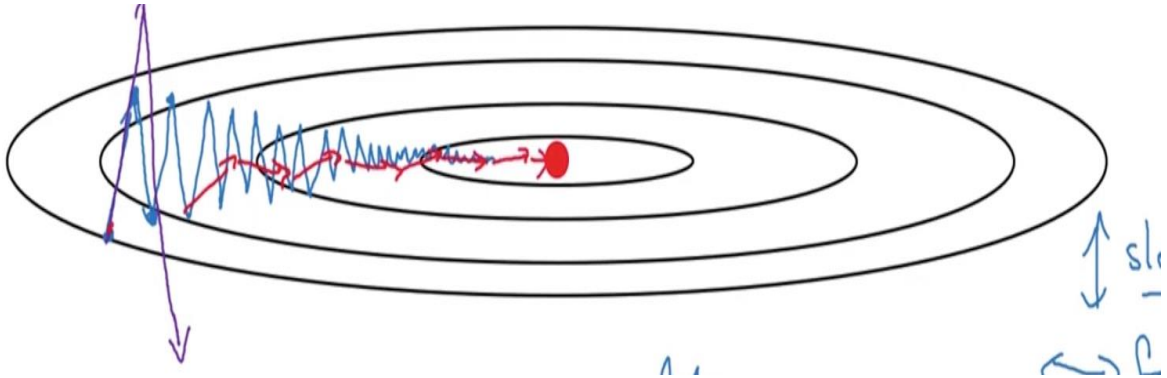

Momentum & RMSprop & Adam

$$Vd_w = \beta_1 * Vd_w + (1-\beta_1) * d_w$$
$$Vd_b = \beta_1 * Vd_b + (1-\beta_1) * d_b$$

} "momentum"-like update

$$Sd_w = \beta_2 * Sd_w + (1-\beta_2) * d_w^{\,2}$$
$$Sd_b = \beta_2 * Sd_b + (1-\beta_2) * d_b^{\,2}$$

} "RMSprop"

$$V_{d_w}^{corrected} = Vd_w \,/\, (1-\beta_1^t)\ , \qquad V_{d_b}^{corrected} = Vd_b \,/\, (1-\beta_1^t)$$
$$S_{d_w}^{corrected} = Sd_w \,/\, (1-\beta_2^t)\ , \qquad S_{d_b}^{corrected} = Sd_b \,/\, (1-\beta_2^t)$$

$$w := w - \alpha * \frac{V_{d_w}^{corrected}}{\sqrt{S_{d_w}^{corrected}} + \varepsilon}$$
$$b := b - \alpha * \frac{V_{d_b}^{corrected}}{\sqrt{S_{d_b}^{corrected}} + \varepsilon}$$

Exponential Weight: $\beta 1, \beta 2, \varepsilon$

Momentum & RMSprop & Adam

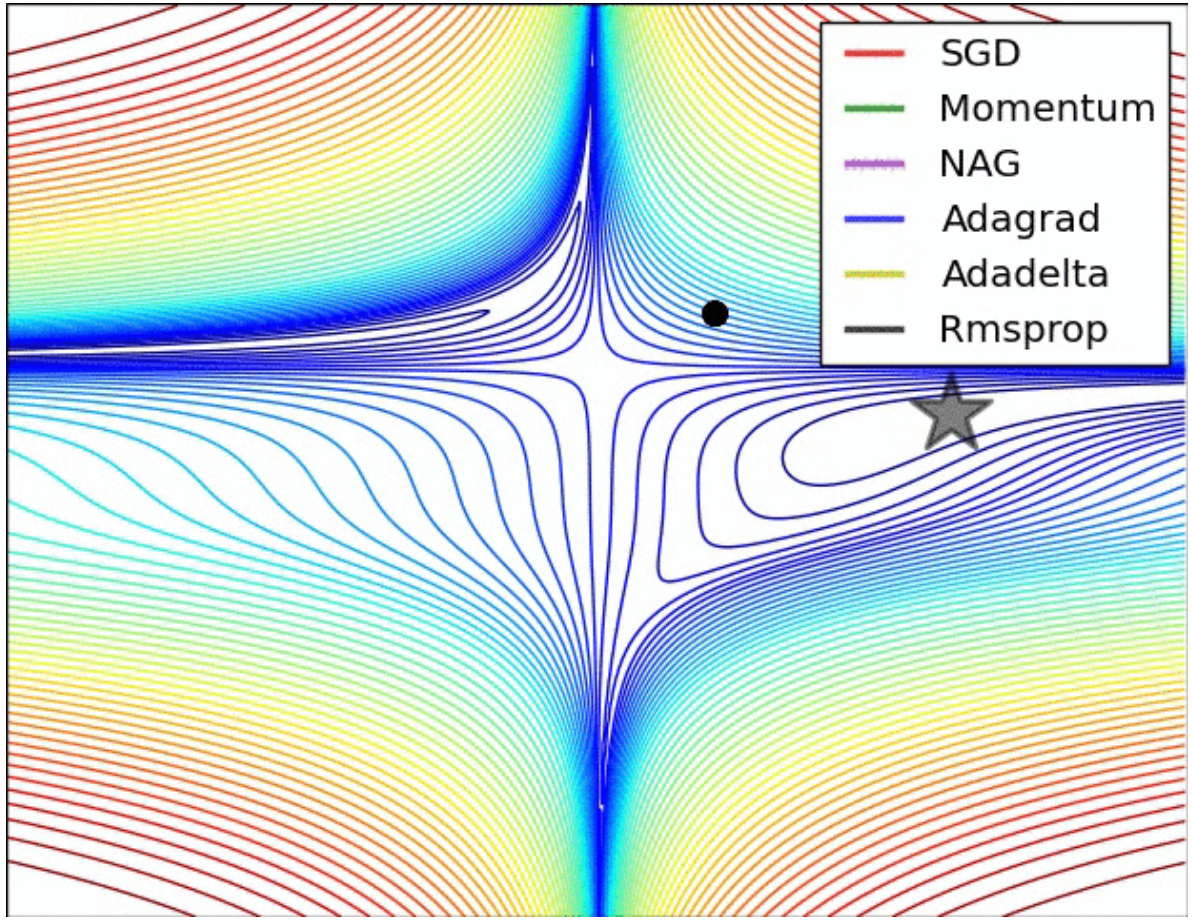

Learning Rate Decay

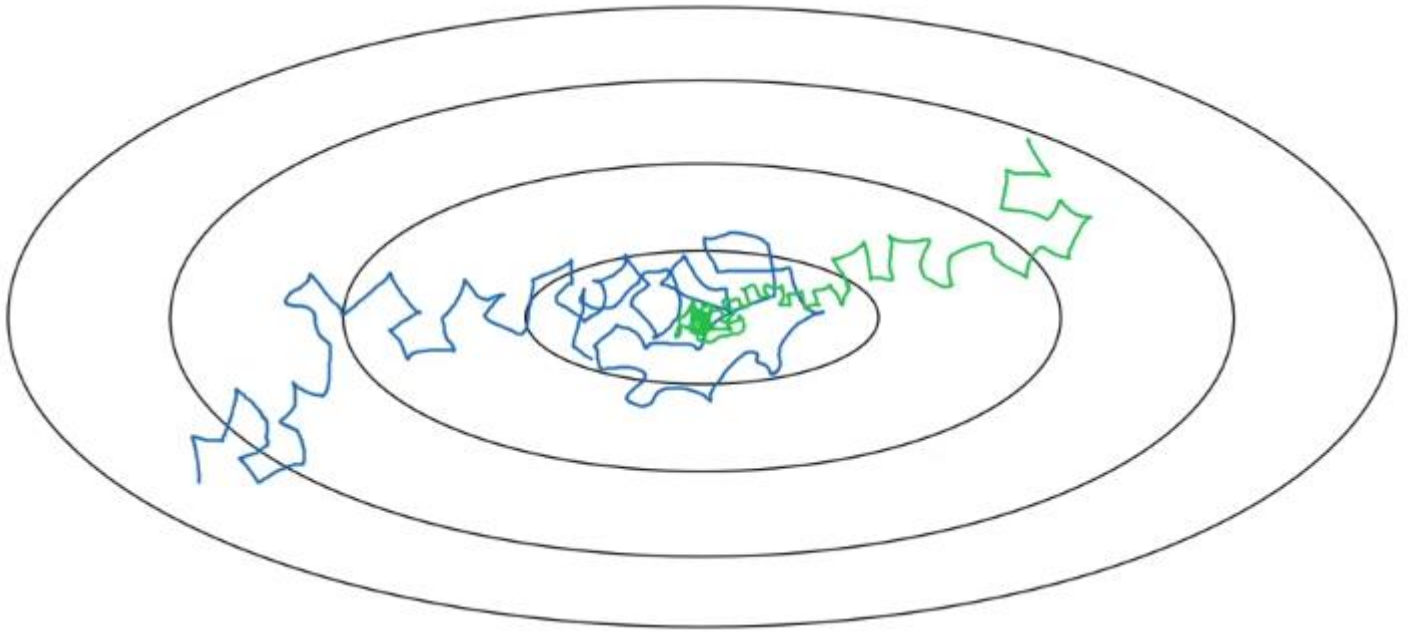

Without Decay

With Decay

Learning Rate Decay: $\alpha = f(\alpha, epoch)$

Assignments

- Gradient Descent
- Mini-Batch Gradient Descent
- Momentum
- Adam
- Model with different optimization algorithms

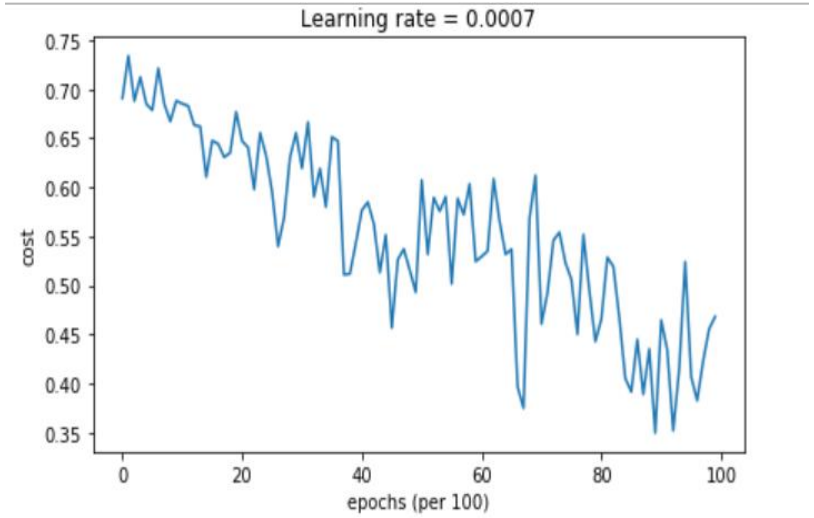

Accuracy: 0.796666666667

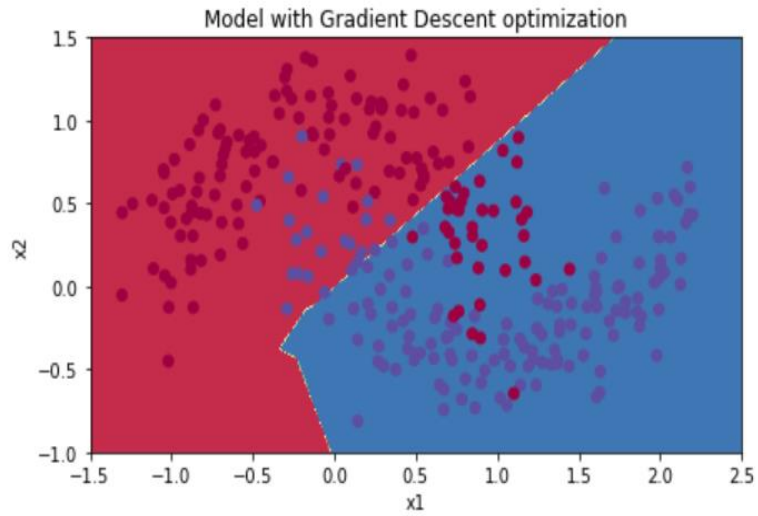

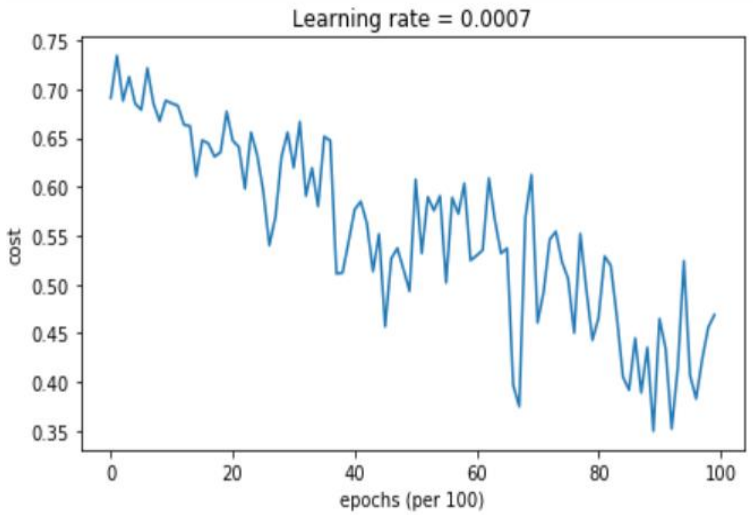

ccuracy: 0.796666666667

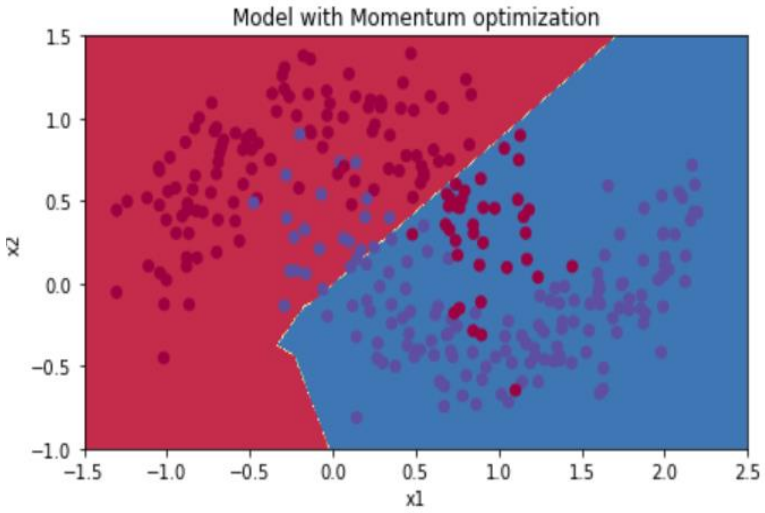

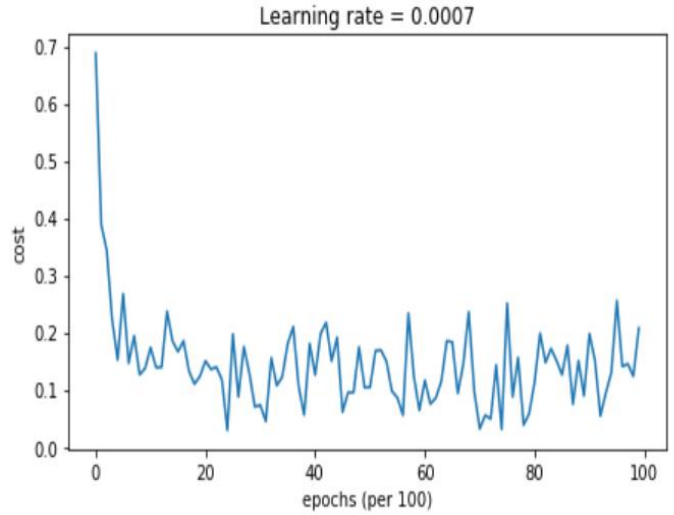

Accuracy: 0.94